



# RC Days 2023

## Advanced topics in HPC:

### *Numerical accuracy, lies and statistics*

Dr. Frank T Willmore  
Boise State University Research Computing  
[FrankWillmore@boisestate.edu](mailto:FrankWillmore@boisestate.edu)

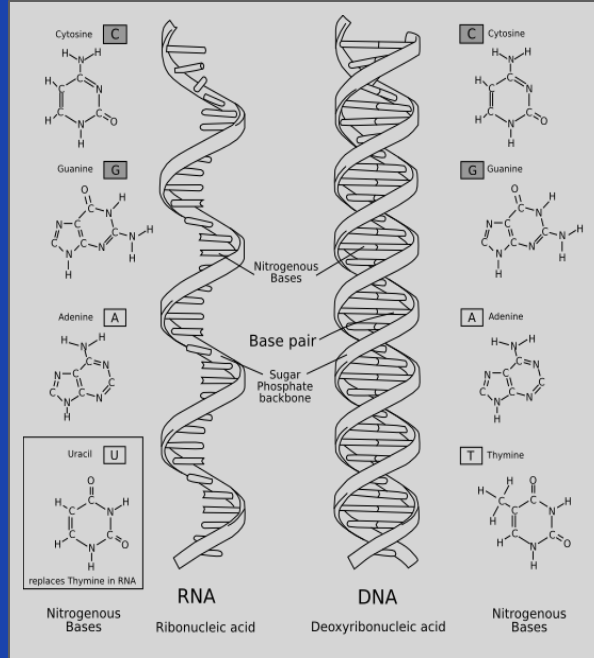


BOISE STATE UNIVERSITY

## What is a number?

- Real
- Imaginary
- Complex
- Integer
- Countable
- Random
- Finite/Infinite
- Positive/Negative
- Rational/Irrational/Transcendental
- Binary/Decimal/Octal/Hexadecimal
- Exact/Approximate
- Floating point

# What is a number? Storage and representation



[https://en.wikipedia.org/wiki/Data\\_storage](https://en.wikipedia.org/wiki/Data_storage)



0x0	0	0	1	1	0	0	0	0
0x1	0	0	0	0	0	0	0	0
0x2	0	0	0	0	0	1	0	0
0x3	0	0	0	0	0	0	0	0
0x4	0	0	1	0	0	0	0	0
0x5	0	0	0	0	0	0	0	0
0x6	0	0	0	0	0	0	0	0
0x7	0	0	0	0	0	0	0	0
0x8	0	0	0	1	0	0	0	0
0x9	0	0	0	0	0	0	0	0
0xA	0	0	1	0	0	0	0	0
0xB	0	0	0	1	0	0	0	0
0xC	0	1	0	0	0	0	0	0
0xD	0	0	0	0	1	0	0	0
0xE	0	0	0	1	0	1	0	0
0xF	0	1	0	0	0	1	0	0




# What is a number? Storage and representation




0x0	0	0	1	1	0	0	0	0
0x1	0	0	0	0	0	0	0	0
0x2	0	0	0	0	0	1	0	0
	0	0	0	0	0	0	0	0
	0	0	1	0	0	0	0	0
	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0
0x8	0	0	0	1	0	0	0	0
0x9	0	0	0	0	0	0	0	0
0xA	0	0	1	0	0	0	0	0
0xB	0	0	0	1	0	0	0	0
0xC	0	1	0	0	0	0	0	0
0xD	0	0	0	0	1	0	0	0
0xE	0	0	0	1	0	1	0	0
0xF	0	1	0	0	0	1	0	0

## What is a number? Machine representation of a floating point number




# What is a number? Machine representation of a floating point number



$-1^{\text{sign}}$   
 $= 1$


# What is a number? Machine representation of a floating point number



$$-1^{\text{sign}} = 1$$

The exponent is 'offset binary' format with a shift of -127, i.e. the exponent is  $(0 \cdot 128) + (1 \cdot 64) + (1 \cdot 32) + (1 \cdot 16) + (1 \cdot 8) + (1 \cdot 4) + (0 \cdot 2) + (0 \cdot 1) - 127$ , or **-3**. This means a factor of  $2^{-3}$ .

# What is a number? Machine representation of a floating point number




$$-1^{\text{sign}} = 1$$

The exponent is 'offset binary' format with a shift of -127, i.e. the exponent is  $(0 \cdot 128) + (1 \cdot 64) + (1 \cdot 32) + (1 \cdot 16) + (1 \cdot 8) + (1 \cdot 4) + (0 \cdot 2) + (0 \cdot 1) - 127$ , or **-3**. This means a factor of  **$2^{-3}$** .

Just like base 10 scientific notation, the exponent is shifted so that the first non-zero digit in the fraction is to the left of the decimal point. In binary, this value ( $2^0=1$ ) is implicit and not recorded. Since only the  $2^{-2}$  bit is set above, the fraction value here is  $2^0 + 2^{-2} = 1 + \frac{1}{4} = \mathbf{1.25}$ .



## What is a number? Machine representation of a floating point number



$$1 \times 2^{-3} \times 1.25 = 0.15625$$

## So what does it all matter? How does this affect my research?

- Machine precision is inherently limited
- This means rounding errors
- Rounding errors are a source of experimental error (in addition to modeling shortfalls) for computational science and can affect results
- Good research accounts for measurement error



## So what does it all matter? How does this affect my research?

- Machine precision is inherently limited
- This means rounding errors
- Rounding errors are a source of experimental error (in addition to modeling shortfalls) for computational science and can affect results
- Good research accounts for measurement error

## What can I do?

- Recognize the limitations of machine precision
- Make responsible and efficient choices regarding precision
  - AI/ML applications often use very low precision
  - GROMACS uses single precision by default
- Document choices
- Do better research

## EXAMPLE: SUMMATION AND ROUNDING

$$1 + x \rightarrow 1 + x; x > \delta$$

$$1 + x \rightarrow 1; x \leq \delta$$

$$y + x \rightarrow \max(x, y); x/y > \delta$$

$$y + x \rightarrow x + y; x/y \leq \delta$$

# EXERCISE: summing floats

```
#include <stdio.h>

int main()
{
    float ttf = 16777216;    // 2^24
    float delta = 1.0f/ttf;  // 0.000000059604645
    float sum = 0.0f;

    // check inputs
    printf("Starting with sum=%1.18f, delta=%1.18f:\n\n", sum, delta);

    // should add to one
    for (int i=0; i<ttf; i++) sum += delta;
    printf("sum=%1.18f\n", sum);

    // what happens if we keep adding?
    for (int i=0; i<ttf; i++) sum += delta;
    printf("sum=%1.18f\n", sum);
}
```



# EXERCISE: summing floats

```
id@borah-login$ dev-session-bsu
...
id@cpulxx$ module load borah-base openmpi/4.1.3/gcc/12.1.0
id@cpulxx$ git clone https://github.com/bsurc/AdvancedHPC
id@cpulxx$ cd AdvancedHPC/machine_numbers_and_MPI
id@cpulxx$ vi add_float.c
id@cpulxx$ make add_float
id@cpulxx$ ./add_float
Starting with sum=0.000000000000000000, delta=0.000000059604644775:

sum=1.000000000000000000
sum=1.000000000000000000
```

# EXERCISE: summing floats

- What happens if you increase the size of ttf?
  - Edit add\_float.c and recompile
  - What is the result?
- What happens if you decrease the size of ttf?
  - Edit add\_float.c and attempt to recompile
  - What does the compiler say?
  - What does it mean?

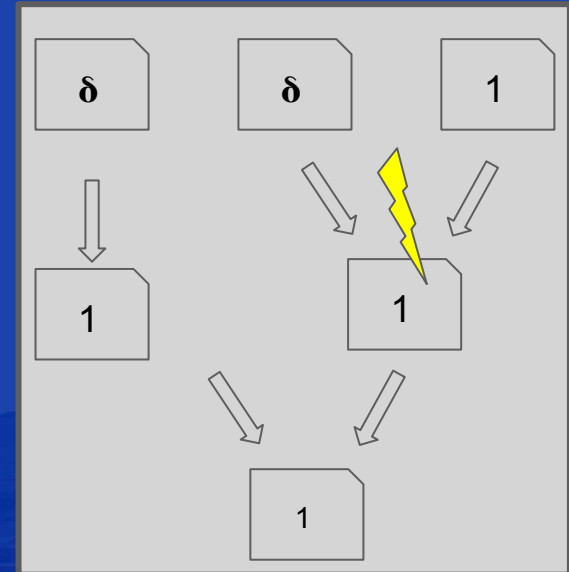
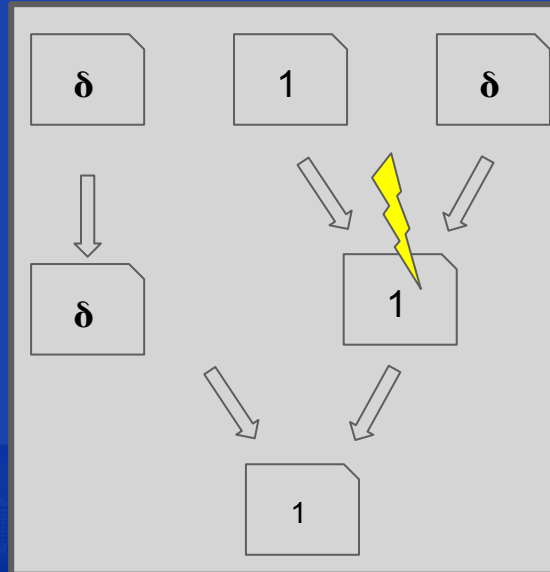

## Example of MPI Reduction

- MPI parallel processes run and complete in an indeterminate order
- Because of rounding, results of a reduction depend on the order in which operations are performed
- MPI operations can give results that vary between runs and depend on size of run



## Example of MPI Reduction

- MPI parallel processes run and complete in an indeterminate order
- Because of rounding(⚡), results of a reduction depend on the order in which operations are performed
- MPI operations can give results that vary between runs and depend on size of run



# EXERCISE: MPI reduction

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char* argv[])
{
    MPI_Init(NULL, NULL);

    float delta = 0.000000059604645;
    int mpi_rank, mpi_size;
    MPI_Comm_rank(MPI_COMM_WORLD, &mpi_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &mpi_size);

    for (int i=0; i<mpi_size; i++)
    {
        float global_sum, local_val;
        if (mpi_rank == i) local_val = 1.0f;
        else local_val = delta;
        MPI_Reduce(&local_val, &global_sum, 1, MPI_FLOAT, MPI_SUM, 0, MPI_COMM_WORLD);
        if (!mpi_rank) printf("big val to rank %d, gives global sum = %1.10f\n", i, global_sum);
    }

    MPI_Finalize( );
}
```



# EXERCISE: MPI reduction

```
id@cpulxx$ vi reduce.c
id@cpulxx$ make reduce
id@cpulxx$ mpirun -n 3 ./reduce
rank 1/3 running.
rank 2/3 running.
rank 0/3 running.
big val to rank 0, gives global sum = 1.0000001192
big val to rank 1, gives global sum = 1.0000000000
big val to rank 2, gives global sum = 1.0000000000
```

# EXERCISE: MPI reduction

- What happens if you run it again?
  - Do the processes still return in the same order?
  - What is the result?
- What happens if you increase the number of processes, e.g. to 4?
- What happens if you load a different MPI (e.g. mpich) and build/run?
  - Do OpenMPI and MPICH give the same results?
  - Why or why not?

## Central Limit Theorem:

The average values of sets of samples of a randomly distributed variable will tend to be distributed normally:

$\mathcal{R} \in [0,1)$  ; for a uniformly distributed random variable  $\mathcal{R}$   
 $R_j \triangleq (1/N) \sum_i \mathcal{R}_i$ ; with average values  $R_j$  for sets of  $N$  samples of  $\mathcal{R}$   
 $p(R_j) \propto \exp\{-(R_j - \langle R_j \rangle)^2/2\}$  ; those average values  $R_j$  are distributed normally

## Central Limit Theorem example:

$\mathcal{R} \in [0,1)$  ; for a uniformly distributed random variable  $\mathcal{R}$

$R_j \triangleq (1/N) \sum_i \mathcal{R}_i$ ; with average values for a sets of N samples of  $\mathcal{R}$

$p(R_j) \propto \exp\{-(R_j - \langle R_j \rangle)^2/2\}$  ; those average values are distributed normally

The experiment:

- Using an MPI code with different random number seeds scattered to different MPI processes, sample values of  $R_j$  are generated and gathered.
- MPI parallel processes then run and complete, although in an indeterminate order
- Because of rounding, results of a reduction depend on the order in which operations are performed
- MPI operations can give results that vary between runs and depend on size of run

# EXERCISE: central limit theorem

```
#include <stdio.h>
#include <mpi.h>
#include <math.h>
#include <stdlib.h>

int main(int argc, char* argv[])
{
    MPI_Init(NULL, NULL);

    int elements_per_proc = 8;
    int initial_seed = 123456;

    int mpi_rank, mpi_size;
    MPI_Comm_rank(MPI_COMM_WORLD, &mpi_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &mpi_size);

    if (!mpi_rank) printf("running %d elements per task with %d tasks.\n", elements_per_proc, mpi_size);

    // create seeds
    int scatter_seeds[mpi_size];
    scatter_seeds[0] = initial_seed;
    if (!mpi_rank) for (int i=1; i<mpi_size; i++) scatter_seeds[i] = scatter_seeds[0] + i;

    // scatter the seeds
    int seed;
    MPI_Scatter(scatter_seeds, 1, MPI_INT, &seed, 1, MPI_INT, 0, MPI_COMM_WORLD);
```



# EXERCISE: central limit theorem

```
// seed the RNG
srand(seed);

float sum = 0;
for (int i=0; i<elements_per_proc; i++)
{
    float sample = (float)rand() / (float)RAND_MAX;
    sum += sample;
}
sum /= elements_per_proc;

// gather and bin the results:
float results[mpi_size];
MPI_Gather(&sum, 1, MPI_FLOAT, results, 1, MPI_FLOAT, 0, MPI_COMM_WORLD);
```

# EXERCISE: central limit theorem

```
int bins[10];
for (int i=0; i<10; i++) bins[i] = 0;

if (!mpi_rank) for (int i=0; i<mpi_size; i++)
{
    printf("binning %f\n", results[i]);
    bins[(int)(results[i]*10)]++;
}

if (!mpi_rank) for (int i=0; i<10; i++)
{
    for (int j=0; j<bins[i]; j++) printf("X");
    printf("\n");
}

MPI_Finalize();
}
```

# EXERCISE: central limit theorem

```
id@cpulxx$ vi central_limit.c
id@cpulxx$ make central_limit
id@cpulxx$ mpirun -n 48 ./central_limit
running 8 elements per task with 48 tasks.
binning 0.488702
binning 0.511513
. . .

XX
XX
XXXXXXXXXXXXXXXXXX
XXXXXXXXXXXXXXXXXXXXXXXXXXXX
XXXXXXX
XX
X
```

# EXERCISE: central limit theorem

- What are the likely sources of error and uncertainty?
- What happens if you run it again for a different seed value?
- What happens if you increase the number of MPI processes?
- What happens if you use a different number of samples for each  $R_j$ ?
- What happens if we have very small values of  $\mathcal{R}$ ?

## Takeaways:

- Greater precision isn't always 'better'
- Knowing when you can use a smaller type can often mean better performance
- Making good choices here depends on understanding the limits of machine numbers
- Operations can give results that vary between runs and depend on size of run
- 'Randomness' is not always an implicit source of error, it depends on how it's used.
- There is a new floating point standard in play (posits) which looks to gain traction
- Limits of machine precision are inherent, no matter the standard

